

BLG 454E Learning From Data

FALL 2022-2023

Assoc. Prof. Yusuf Yaslan

These Slides are adopted by

Assoc. Prof. Yusuf Yaslan & Assist. Prof. Ayşe Tosun

Dimensionality Reduction

Why Reduce Dimensionality?

- Ideally, we should not need feature selection or extraction as a separate process; the classifier (or regressor) should be able to use whichever features are necessary, discarding the irrelevant.
- There are several reasons why we are interested in reducing dimensionality as a separate preprocessing step:

Why Reduce Dimensionality?


- Reduces time complexity: Less computation
- Reduces space complexity: Less parameters
- Saves the cost of observing the feature
- Simpler models are more robust on small datasets
- More interpretable; simpler explanation
- Data visualization (structure, groups, outliers, etc) if plotted in 2 or 3 dimensions

Curse of Dimensionality

- A term coined by Bellman in 1961
- Refers to the problems associated with multivariate data analysis as the dimensionality increases
- Describes the increasing difficulty in training a model when more predictor variables are added to it
- For a given sample size, there is a maximum number of features above which the performance of our classifier will degrade rather than improve

Curse of Dimensionality

- In most cases, the additional information that is lost by discarding some features is (more than) compensated by a more accurate mapping in the lower-dimensional space



Feature Selection vs Extraction

- **Feature selection:** Choosing $k < d$ important features, ignoring the remaining $d - k$
Subset selection algorithms
- **Feature extraction:** Project the original $x_i, i = 1, \dots, d$ dimensions to new $k < d$ dimensions, $z_j, j = 1, \dots, k$

Principal components analysis (PCA), linear discriminant analysis (LDA), factor analysis (FA)

Subset Selection

- We are interested in finding the best subset of the set of features. The best subset contains the least number of dimensions that most contribute to accuracy. We discard the remaining, unimportant dimensions.
- There are 2^d subsets of d features but we cannot test for all of them unless d is small and we employ heuristics to get a reasonable (but not optimal) solution in reasonable (polynomial) time.

Subset Selection: Search Strategy and Objective Function

- A **search strategy** is used to select candidate subsets. i.e sequential, randomized and exponential algorithms
- An **objective function** is used to evaluate these candidates
- Objective functions are divided in two groups
- **Filters**: evaluate subsets by their information content, e.g., interclass distance, statistical dependence or information-theoretic measures
- **Wrappers**: use a classifier to evaluate subsets by their predictive accuracy (on validation data) by statistical resampling or cross-validation

Subset Selection


- **Forward search:** Add the best feature at each step
 - Set of features F initially \emptyset .
 - At each iteration, find the best new feature
 $j = \operatorname{argmin}_i E(F \cup x_i)$ % $E()$: Error on the validation set
 - Add x_j to F if $E(F \cup x_j) < E(F)$
- Hill-climbing $O(d^2)$ algorithm (train and test the system $d, d-1, d-2, \dots, (d-k)$ times)

Subset Selection

- This is a local search procedure and does not guarantee finding the optimal subset, namely, the minimal subset causing the smallest error
- **Backward search:** Start with all features and remove one at a time, if possible.
- Floating search (Add k , remove l) number of added features and removed features can also change at each step.

Principal Components Analysis (PCA)


- Find a low-dimensional space such that when \mathbf{x} is projected there, information loss is minimized.



Vector Projections

- The red vector has length $\text{length} = \|\mathbf{u}\| \cos \theta = \frac{\mathbf{u} \cdot \mathbf{v}}{\|\mathbf{v}\|}$,
- $\mathbf{v}/\|\mathbf{v}\|$ is a unit vector in the direction of \mathbf{v} .
- The dashed red vector

$$\text{dashed red vector} = \left(\frac{\mathbf{u} \cdot \mathbf{v}}{\mathbf{v} \cdot \mathbf{v}} \right) \mathbf{v}.$$



Principal Components Analysis (PCA)


- Find a low-dimensional space such that when \mathbf{x} is projected there, information loss is minimized.
- The projection of \mathbf{x} on the direction of \mathbf{w} is: $z = \mathbf{w}^T \mathbf{x}$
- Find \mathbf{w} such that $\text{Var}(z)$ is maximized

$$\begin{aligned}\text{Var}(z) &= \text{Var}(\mathbf{w}^T \mathbf{x}) = E[(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \boldsymbol{\mu})^2] \\ &= E[(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \boldsymbol{\mu})(\mathbf{w}^T \mathbf{x} - \mathbf{w}^T \boldsymbol{\mu})] \\ &= E[\mathbf{w}^T (\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T \mathbf{w}] \\ &= \mathbf{w}^T E[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T] \mathbf{w} = \mathbf{w}^T \Sigma \mathbf{w}\end{aligned}$$

where $\text{Var}(\mathbf{x}) = E[(\mathbf{x} - \boldsymbol{\mu})(\mathbf{x} - \boldsymbol{\mu})^T] = \Sigma$

Principal Components Analysis (PCA)

- Maximize $\text{Var}(z)$ subject to $\|\mathbf{w}\|=1$
- Use Lagrange Multipliers:
- $\max f(x,y)$ subject to $g(x,y)=c$ we need both f and g continuous with first partial derivative.
- We introduce a new variable (λ) Lagrange Multiplier and optimize:
- $L(x,y, \lambda) = f(x,y) - \lambda(g(x,y)-c)$



- Maximize $\text{Var}(z)$ subject to $\|\mathbf{w}\|=1$

$$\max_{\mathbf{w}_1} \mathbf{w}_1^T \Sigma \mathbf{w}_1 - \alpha (\mathbf{w}_1^T \mathbf{w}_1 - 1)$$

Take the derivative w.r.t. \mathbf{w}_1 and set equal to 0.

$$2\Sigma \mathbf{w}_1 - 2\alpha \mathbf{w}_1 = 0$$

$\Sigma \mathbf{w}_1 = \alpha \mathbf{w}_1$ that is, \mathbf{w}_1 is an eigenvector of Σ

Choose the one with the largest eigenvalue for $\text{Var}(z)$ to be max

- Second principal component:
- Max $\text{Var}(z_2)$, s.t., $\|\mathbf{w}_2\|=1$ and orthogonal to \mathbf{w}_1
- We should add the second constraint with β Lagrange multiplier

$$\max_{\mathbf{w}_2} \mathbf{w}_2^T \Sigma \mathbf{w}_2 - \alpha (\mathbf{w}_2^T \mathbf{w}_2 - 1) - \beta (\mathbf{w}_2^T \mathbf{w}_1 - 0)$$

- Take the derivative w.r.t. \mathbf{w}_2 and set equal to 0.

$$2\Sigma \mathbf{w}_2 - 2\alpha \mathbf{w}_2 - \beta \mathbf{w}_1 = 0$$

- Premultiply with \mathbf{w}_1^T

$$2\mathbf{w}_1^T \Sigma \mathbf{w}_2 - 2\alpha \mathbf{w}_1^T \mathbf{w}_2 - \beta \mathbf{w}_1^T \mathbf{w}_1 = 0 \qquad 2\alpha \mathbf{w}_1^T \mathbf{w}_2 = 0$$

$$\mathbf{w}_1^T \Sigma \mathbf{w}_2 = \mathbf{w}_2^T \Sigma \mathbf{w}_1 \qquad \Sigma \mathbf{w}_1 = \lambda_1 \mathbf{w}_1 \qquad \lambda_1 \mathbf{w}_2^T \mathbf{w}_1 = 0$$

- Then $\beta = 0$


$\Sigma \mathbf{w}_2 = \alpha \mathbf{w}_2$ that is, \mathbf{w}_2 is another eigenvector of Σ and so on.

What PCA does

$$\mathbf{z} = \mathbf{W}^T(\mathbf{x} - \mathbf{m})$$

where the columns of \mathbf{W} are the eigenvectors of Σ , and \mathbf{m} is sample mean

Centers the data at the origin and rotates the axes



How to choose k ?

- Proportion of Variance (PoV) explained

$$\frac{\lambda_1 + \lambda_2 + \dots + \lambda_k}{\lambda_1 + \lambda_2 + \dots + \lambda_k + \dots + \lambda_d}$$

when λ_i are sorted in descending order

- Typically, stop at $\text{PoV} > 0.9$
- Scree graph plots of PoV vs k , stop at “elbow”

